

Enzyme Classification Using ProtBERT

Rime Tarchihi-M2 ACSYON EUR

December 17, 2025

Dataset Collection

UniProt BLAST Align Peptide search ID mapping SPARQL UniProtKB annotation Advanced | List Search

Status

- All reviewed (Swiss-Prot) entries
- Selected unreviewed (TrEMBL) entries with experimental or biologically important data
- Entries to be removed: Unreviewed (TrEMBL) entries that are not part of a reference proteome

Entries removed from unreviewed UniProtKB/TrEMBL will remain accessible in the UniParc sequence archive. Please read our [help page](#), view [affected entries](#) and [proteomes](#), or [contact us](#) with any questions.

UniProtKB 20,226 results

Tools Download (20k) Add View: Cards Table Customize columns Share

Group by Enzyme Classification

Search for Enzyme Classification

Enter Enzyme Classification name or ID

UniProtKB Entries	Enzyme Classification
20,226	Top level

Group by

- Taxonomy
- Keywords
- Gene Ontology
- Enzyme Class

Proteins with

Hydrolases 30%
Isomerases 4%
Ligases 5%
Lyases 6%
Oxidoreductases 16%
Transferases 36%
Translocases 2%

Environment Setup and Data Loading

```
[ ] !pip install -q transformers datasets evaluate scikit-learn accelerate sentencepiece
[ ]                                     84.1/84.1 kB 3.4 MB/s eta 0:00:00
[ ]
import pandas as pd
df = pd.read_csv(
    "uniprotkb_reviewed_ec.tsv.gz",
    sep="\t",
    compression="gzip"
)
```

Dataset Exploration

- I inspected the dataset structure and column names
- I displayed the first few protein entries to understand the data format

```
[ ] print(df.head())
print(df.columns)
print("Total sequences:", len(df))

Entry      Entry Name          Sequence \
0 A0A0B4K692 NEP2_DROME MQTVIQNPNWRRRNKLEKSLLVSLGIMFWLATGFLWIGKVLRT...
1 A0A0B4K7J2 RBP2_DROME MFTTRKEVDAHVKMLGKLQPGRERDIKGLAVARLYMKVQEYPKAI...
2 A0A0B4KEE4 KOI_DROME MSENTYQIETRRRSRSKTPFLRSSCDHENCEHAGEEGHVHHLKRKS...
3 A0A0B4KGY6 NOVA_DROME MESIMKVAMDAAEQLIQQFGFDYLQQQLQLQHQNQHNSSPQQPQH...
4 A0A0B4LFY9 STING_DROME MAIASNVVEAGNAVRAEKGRKYFYFRKMIQDYZIDTSIRIVATVFLA...

EC number
0 3.4.24.11
1 2.3.2.-
2 NaN
3 NaN
4 NaN
Index(['Entry', 'Entry Name', 'Sequence', 'EC number'], dtype='object')
Total sequences: 20226
```

Removing Non-Enzyme Proteins

- I removed protein entries that do not have an EC number
- Proteins without EC annotations are not enzymes and cannot be used for classification
- This step ensures that all remaining samples belong to valid enzyme classes

[]



```
# Remove proteins without EC numbers
df = df.dropna(subset=["EC number"])

print("After removing non-enzymes:", len(df))
```

▼

... After removing non-enzymes: 7441

EC Class Extraction and Label Definition

```
# If multiple EC numbers exist, keep the first
df["EC_number"] = df["EC number"].apply(lambda x: x.split(";")[0])

# Extract first EC digit
df["ec_class"] = df["EC_number"].apply(lambda x: x.split(".")[0])

# Keep only valid enzyme classes 1-7
df = df[df["ec_class"].isin([str(i) for i in range(1, 8)])]

print(df["ec_class"].value_counts())
```

```
ec_class
2    2705
3    2225
1    1239
4     439
6     383
5     289
7     161
Name: count, dtype: int64
```

Label Mapping and Final Dataset Format

```
▶ ec_map = {  
    "1": "Oxidoreductases",  
    "2": "Transferases",  
    "3": "Hydrolases",  
    "4": "Lyases",  
    "5": "Isomerases",  
    "6": "Ligases",  
    "7": "Translocases"  
}  
  
df["label_name"] = df["ec_class"].map(ec_map)  
df["label"] = df["ec_class"].astype(int) - 1 # 0-6  
  
df = df[["Sequence", "label", "label_name"]]  
df.head()
```

...

	Sequence	label	label_name	grid icon	row icon
0	MQTVIQNPNNWRRRNKLEKSLLVSLGIMFVVLATGFLWIGKVLRT...	2	Hydrolases		
1	MFTTRKEVDAHVHJKLMQLQPGRERDIKGHLAVARLYMKVQEYPKAI...	1	Transferases		
6	MNLTKLMKVFGYINIITNCVQSFTNRADKKRYNVFAKSFINTINTN...	2	Hydrolases		
7	MPPRCRRLPLLFILLLAVRPLSAAAASSIAAAPASSYRRISWASNLI...	1	Transferases		
8	MASPPPFDICGDLDDDPPTPPAPTPLAAPTPNGLNDRLLRTRTHQR...	2	Hydrolases		

Train, validation, and test Split

```
from sklearn.model_selection import train_test_split
# 70% train, 15% val, 15% test
train_df, temp_df = train_test_split(
    df,
    test_size=0.3,
    stratify=df["label"],
    random_state=42
)

val_df, test_df = train_test_split(
    temp_df,
    test_size=0.5,
    stratify=temp_df["label"],
    random_state=42
)

print("Train size:", len(train_df))
print("Validation size:", len(val_df))
print("Test size:", len(test_df))
```

```
Train size: 5208
Validation size: 1116
Test size: 1117
```

Sequence Preprocessing for ProtBERT

- I preprocessed protein sequences to match ProtBERT input requirements
- I separated each amino acid with a space
- This allows ProtBERT to tokenize amino acids individually
- The same preprocessing was applied to training, validation, and test sets

```
def space_separate(seq):
    return " ".join(list(seq))

train_df["Sequence"] = train_df["Sequence"].apply(space_separate)
val_df["Sequence"] = val_df["Sequence"].apply(space_separate)
test_df["Sequence"] = test_df["Sequence"].apply(space_separate)
```

```

import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification
MODEL_NAME = "Rostlab/prot_bert"
tokenizer = AutoTokenizer.from_pretrained(
    MODEL_NAME,
    do_lower_case=False
)
model = AutoModelForSequenceClassification.from_pretrained(
    MODEL_NAME,
    num_labels=7
)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
print("on", device)

```

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
... /usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as s
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
tokenizer_config.json: 100% ██████████ 86/0/86.0 [00:00<00:00, 4.01kB/s]
config.json: 100% ██████████ 361/361 [00:00<00:00, 14.2kB/s]
vocab.txt: 100% ██████████ 81.0/81.0 [00:00<00:00, 2.32kB/s]
special_tokens_map.json: 100% ██████████ 112/112 [00:00<00:00, 4.62kB/s]
pytorch_model.bin: 100% ██████████ 1.68G/1.68G [00:27<00:00, 36.7MB/s]
model.safetensors: 20% ████████ 343M/1.68G [00:10<00:24, 54.0MB/s]
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at Rostlab/prot_bert and are newly init
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
on cuda

```

Sequence Tokenization

- I defined a custom tokenization function for protein sequences.
- Each amino acid sequence is converted into ProtBERT tokens.
- Sequences are padded to a fixed length of 512 tokens.
- Longer sequences are truncated to fit the model input size.

```
def tokenize(batch):
    return tokenizer(
        batch["Sequence"],
        padding="max_length",
        truncation=True,
        max_length=512
    )
```

Dataset Preparation

- I converted the training and validation dataframes into Hugging Face Dataset objects.
- The tokenization function was applied in batch mode for efficiency.
- Non-essential columns (Sequence and label_name) were removed after tokenization.

```
from datasets import Dataset
train_dataset = Dataset.from_pandas(train_df)
val_dataset = Dataset.from_pandas(val_df)
train_dataset = train_dataset.map(tokenize, batched=True)
val_dataset = val_dataset.map(tokenize, batched=True)
train_dataset = train_dataset.remove_columns(["Sequence", "label_name"])
val_dataset = val_dataset.remove_columns(["Sequence", "label_name"])
train_dataset.set_format("torch")
val_dataset.set_format("torch")
```

Map: 100%  5208/5208 [00:06<00:00, 908.12 examples/s]

Map: 100%  1116/1116 [00:01<00:00, 1056.19 examples/s]

Test Dataset Preparation

- I created a separate test dataset from the test split.
- The same tokenization process was applied to ensure consistency with training data.
- Original text columns were removed after tokenization.
- The test dataset was formatted as PyTorch tensors.
- This dataset is used only for final, unbiased model evaluation.

```
[15] #test dataset
✓ 3s
  test_dataset = Dataset.from_pandas(test_df)
  test_dataset = test_dataset.map(tokenize, batched=True)
  test_dataset = test_dataset.remove_columns(["Sequence", "label_name"])
  test_dataset.set_format("torch")

  *** Map: 100% [1117/1117 [00:03<00:00, 355.43 examples/s]
```

Training Configuration

```
[16] ✓ 0s   from transformers import TrainingArguments  
  
   training_args = TrainingArguments(  
       output_dir="../protbert_ec",  
       eval_strategy="epoch",  
       save_strategy="epoch",  
       learning_rate=2e-5,  
       per_device_train_batch_size=4,  
       per_device_eval_batch_size=4,  
       num_train_epochs=3,  
       weight_decay=0.01,  
       logging_steps=50,  
       load_best_model_at_end=True,  
       metric_for_best_model="accuracy",  
       fp16=torch.cuda.is_available(),  
       report_to="none"  
   )
```

Evaluation Metrics Definition

```
[17] ✓ 0s
▶ from sklearn.metrics import accuracy_score, precision_recall_fscore_support
import numpy as np
def compute_metrics(eval_pred):
    logits, labels = eval_pred
    preds = np.argmax(logits, axis=1)

    acc = accuracy_score(labels, preds)
    precision, recall, f1, _ = precision_recall_fscore_support(
        labels, preds, average="weighted"
    )
    return {
        "accuracy": acc,
        "precision": precision,
        "recall": recall,
        "f1": f1
    }
```

Here, I define the evaluation metrics used during model training and validation. I compute accuracy, precision, recall, and F1-score using a weighted average to handle class imbalance across enzyme classes.

Trainer Initialization

```
[18] ✓ 2s
from transformers import Trainer

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    tokenizer=tokenizer,
    compute_metrics=compute_metrics
)
/tmp/ipython-input-293716929.py:3: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `Trainer.__init__`
trainer = Trainer(
```

In this step, I initialize the Hugging Face Trainer by connecting the ProtBERT model, training arguments, training and validation datasets, tokenizer, and the evaluation metrics function. This configuration manages the full training and evaluation pipeline.

Training Results Analysis

```
[19] ✓ 44m
  ● trainer.train()
...
[3906/3906 44:01, Epoch 3/3]
Epoch Training Loss Validation Loss Accuracy Precision Recall F1
1 0.874300 0.712859 0.793907 0.779327 0.793907 0.771268
2 0.317900 0.581317 0.861111 0.867967 0.861111 0.858107
3 0.348100 0.494214 0.893369 0.896536 0.893369 0.893761
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and _warn_prf(average, modifier, f"({metric.capitalize()}) is", len(result))
TrainOutput(global_step=3906, training_loss=0.5777749984739258, metrics={'train_runtime': 2643.7921, 'train_samples_per_second': 5.91, 'train_steps_per_second': 1.477, 'total_flos': 1.8188221943291904e+16, 'train_loss': 0.5777749984739258, 'epoch': 3.0})
```

- I observe a consistent decrease in training and validation loss across epochs.
- Classification accuracy improves from **79.4%** in epoch 1 to **89.3%** in epoch 3.
- Precision, recall, and F1-score increase steadily, indicating better class discrimination.
- Validation performance closely follows training performance, suggesting good generalization.
- No strong signs of overfitting are observed within the 3 training epochs.

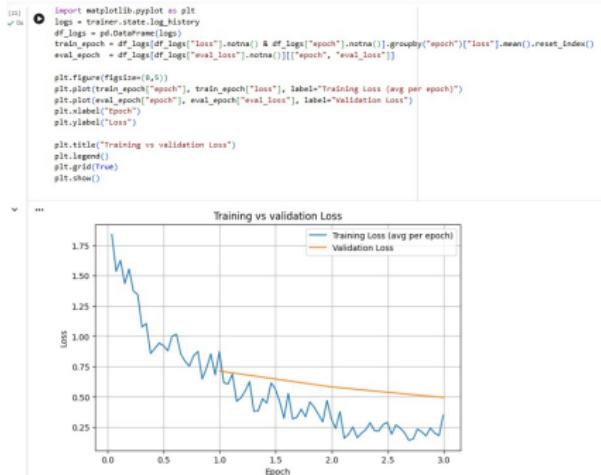
Validation Set Evaluation

```
[20] ✓ 34s trainer.evaluate()  
... [279/279 00:34]  
{'eval_loss': 0.49421370029449463,  
 'eval_accuracy': 0.8933691756272402,  
 'eval_precision': 0.8965362131355358,  
 'eval_recall': 0.8933691756272402,  
 'eval_f1': 0.8937610405614278,  
 'eval_runtime': 34.5459,  
 'eval_samples_per_second': 32.305,  
 'eval_steps_per_second': 8.076,  
 'epoch': 3.0}
```

Validation Results

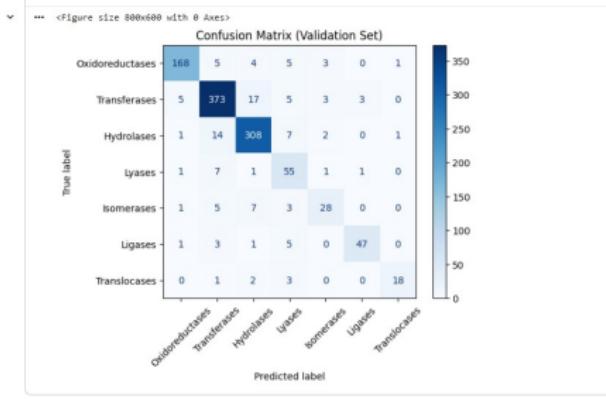
- I evaluated the trained model on the validation set after 3 epochs.
- The model achieves an accuracy of 89.34%.
- Precision, recall, and F1-score are all close to 0.89, indicating balanced performance.
- The low validation loss confirms stable learning and good generalization.
- These results guided the selection of the best model checkpoint.

Training vs Validation Loss

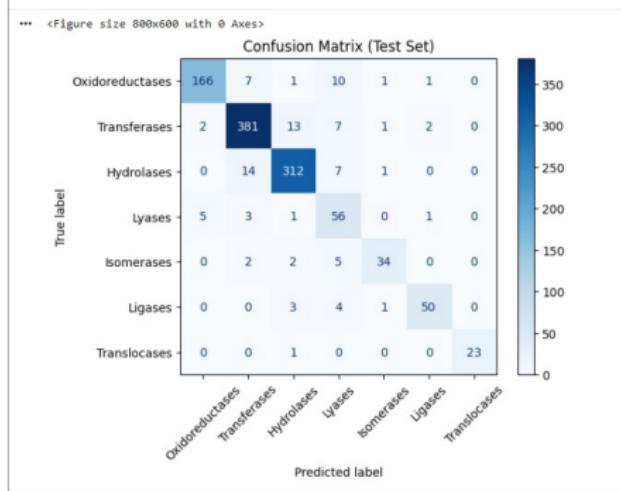


- I monitored both training and validation loss across epochs.
- Training loss decreases steadily.
- Validation loss also decreases and stabilizes over time.
- This indicates good generalization and no strong overfitting.

Confusion Matrices



Validation Set



Test Set

Both confusion matrices show strong diagonal dominance, indicating accurate classification across enzyme classes with limited confusion.

RimeEtudiante / enzyme-classification-protbert

Type to search

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

enzyme-classification-protbert

main 1 Branch 0 Tags

Go to file Add file Code

RimeEtudiante Add files via upload a089710 · now 3 Commits

.gitignore Initial commit 5 minutes ago

DL_Rime (2).pdf Add files via upload now

README.md Initial commit 5 minutes ago

Rime-DeepLearningProjectLipynb Add files via upload now

README

enzyme-classification-protbert

Enzyme classification using ProtBERT and protein sequence data

About

Enzyme classification using ProtBERT and protein sequence data

Readme Activity 0 stars 0 watching 0 forks

Releases

No releases published Create a new release

Packages

No packages published Publish your first package

Languages

References

- Uni24 The UniProt Consortium, *UniProt: the Universal Protein Knowledgebase*, Nucleic Acids Research, 2024.
- EC99 Webb, E. C., *Enzyme Nomenclature*, IUBMB, 1992.
- Cos24 Coste, F., *Enzymatic annotation of protein sequences with a deep language model*, MERIT CNRS Network, 2024.

Thank you for your attention